

Neural and Evolutionary Computation (NEC)

A1: Prediction with Back-Propagation and Linear Regression

Objective

Prediction using the following algorithms:

- **Back-Propagation (BP)**, implemented by the student
- **Multiple Linear Regression (MLR)**, using free software

Datasets

The predictions must be performed on three datasets:

1. File: *AI-turbine.txt*
 - 5 features: the first 4 are the input variables, the last one is the value to predict
 - 451 patterns: use the first 85% for training and validation, and the remaining 15% for test
2. File: *AI-synthetic.txt*
 - 10 features: the first 9 are the input variables, the last one is the value to predict
 - 1000 patterns: use the first 80% for training and validation, and the remaining 20% for test
3. Search a **dataset from the Internet**, with the following characteristics:
 - At least 6 features, one of them used for prediction
 - The prediction variable must take real (float or double) values; it should not represent a categorical value (that would correspond to a classification task)
 - At least 400 patterns
 - Select randomly 80% of the patterns for training and validation, and the remaining 20% for test; it is important to shuffle the original data, to destroy any kind of sorting it could have
 - In case of doubt on the suitability of the selected data, please ask the professor
 - Add to the documentation of this assignment the link to the source webpage

Procedure

- We want to simulate a real application, thus:
 - Apply the techniques of **data preprocessing** to the third dataset, the one you have found on the Internet; the provided datasets are already cleaned, no need to preprocess them. Basically, check for missing values, represent correctly categorical values, and look for outliers
 - **Data normalization** is necessary for BP training of all datasets
- You must find good values for all the **parameters of BP**: architecture of the network, learning rate and momentum, activation function, and number of epochs. Try different combinations from a certain set and use the expected error to select the best-

performing set of parameters. Once you have found those best parameters, apply them to the test set.

- Spend some time trying to **automate the whole process**, otherwise you will spend a lot of time repeating the same manual changes of the parameters

Implementation of BP

- Choose any **programming language**, but with the following restrictions:
 - Julia ≥ 1.8 , Python ≥ 3.6 , R ≥ 4.0 , Matlab $\geq 2020a$, Java ≥ 8
 - Libraries that implement neural networks and BP are **forbidden**
 - Libraries that implement the reading of the data and the scaling of the features are accepted
 - Plotting libraries are also accepted
- It is recommended to use a compiled programming language (Julia, C, C++, C#, Java or Ada) instead of interpreted ones (Python, Matlab, R) because BP requires a lot of calculations, which can take seconds or minutes with the former, and hours with the later ones.
- The **implementation** must be based on the algorithm and equations in **document [G] of Unit 3** at Moodle
- The implementation must use the following **variables** to hold all the information about the structure of the multilayer neural network and the BP:
 - L : number of layers
 - n : an array with the number of units in each layer (including the input and output layers)
 - h : an array of arrays for the fields (h)
 - $\mathbf{x_i}$: an array of arrays for the activations (ξ)
 - w : an array of matrices for the weights (w)
 - θ : an array of arrays for the thresholds (θ)
 - δ : an array of arrays for the propagation of errors (Δ)
 - d_w : an array of matrices for the changes of the weights (δw)
 - d_θ : an array of arrays for the changes of the weights ($\delta \theta$)
 - d_w_{prev} : an array of matrices for the previous changes of the weights, used for the momentum term ($\delta w^{(prev)}$)
 - d_θ_{prev} : an array of arrays for the previous changes of the thresholds, used for the momentum term ($\delta \theta^{(prev)}$)
 - $fact$: the name of the activation function that it will be used. It can be one of these four: sigmoid, relu, linear, tanh.
- For example, the weight $w_{ij}^{(L)}$ between unit j in layer $L-1$ and unit i in layer L is accessed as $w[L][i, j]$ in Julia and C#, or $w[L][i][j]$ in C and Java
- The idea behind this structure is that the code must be able to deal with **arbitrary multilayer networks**. For example, a network with architecture 3:9:5:1 (4 layers, 3 input units, 1 output unit, and two hidden layers with 9 and 5 units, respectively), would have $n=[3; 9; 5; 1]$, and $\mathbf{x_i}$ would be an array of length 4 (one component per layer), with $\mathbf{x_i}[1]$ and array of real numbers of length 3, $\mathbf{x_i}[2]$ and array of real numbers of length 9, $\mathbf{x_i}[3]$ and array of real numbers of length 5, and $\mathbf{x_i}[4]$ and array of real numbers of length 1. Similarly, $w[2]$ would be an array 9×3 , $w[3]$ an array 5×9 , and $w[4]$ and array 1×5 ; $w[1]$ is not used.

- Additionally, the use of this structure, name conventions and array dimensions makes it easy to convert the equations into code

Training parameters and execution

- All the **training parameters** must be put in a **text file**. It must include:
 - Name of the data file
 - Number of training and test patterns
 - Information about the type of activation function
 - Number of layers
 - Number of units in each layer
 - Number of epochs
 - Learning rate and momentum
 - Optionally: information about the scaling method (normalization or standardization) of inputs and/or outputs, and in the case of normalization, the range of the normalized data
 - Optionally: the selected activation function (sigmoid, tanh, ReLU, etc.)
 - Optionally: name of output file(s)
- If compilation is needed, add a **script for compilation**, e.g., a file `compile.sh`, `compile.bat`, or a Makefile
- The **execution** should be something like this:
 - `./run_bp parameters_file.txt`
 - `./run_bp.exe parameters_file.txt`
 - `./run_bp.sh parameters_file.txt`
- Comment: for the search of the best-performing parameters you may need to work with other scripts and/or programs; that's acceptable, but the availability of the execution of BP using a parameters file is mandatory, even if this is not what you have used to obtain your results
- **Plots** of training and validation quadratic errors as a function of the epoch are useful to decide the number of epochs and to avoid over-fitting (over-training)

Multilinear regression

- There is no need to implement MLR, you may use any free tool you want. Anyway, I recommend using libraries from Julia, Python, R, Octave or Matlab, since they allow to use them with a simple script
- Other alternatives, like Weka, Excel and others require the use of a graphical interface, thus they are not easily automated and we don't recommend its usage

Evaluation of the results

- Although BP and MLR are based on the minimization of the mean squared error, it does not constitute a useful measure of the prediction error. We will use the **mean absolute percentage error (MAPE)**, given by:

$$E(\%) = 100 \frac{\sum_{\mu} |y^{\mu} - z^{\mu}|}{\sum_{\mu} z^{\mu}}$$

- Evaluate the previous measure (MAPE) on the Test and Validation sets for both BP and MLR, and for all the datasets
- The best way to visualize the results is with **scatter plots** of the prediction value y^{μ} compared with the real value z^{μ} . The closer the points are to the diagonal, the better the prediction. You must include them in the documentation

Delivery

- This assignment can be done **alone or in pairs** (groups of two)
- The **delivery** must be done in a compressed file (e.g., zip, rar, tgz) whose name should be of the form:
 - A1-Name1Surname1-Name2Surname2.zip
- The **delivery** must contain:
 - **Files:**
 - Source code of BP (avoiding intermediate files, e.g., *.o, *.class, etc)
 - Source code of MLR, if any
 - Parameters files (for the best-performing BP)
 - Script files
 - Results files for MLR and for your best-performing BP: CSV (or TXT tab separated) files with 2 columns, the real value z^{μ} and the prediction y^{μ} for the test sets
 - Other result files (e.g., weights and threshold of the best-performing trained networks, etc.)
 - **Report** in pdf:
 - Description of the implementation (languages, tools used, etc.)
 - Execution instructions
 - Implementation decisions
 - Description and link to the selected dataset
 - Predictions
 - Evaluation of the predictions: test error
 - Scatter plots of the prediction versus real value for both BP and MLR on the Test subsets
 - Discussion and interpretation of the results